

akash-adease-time-series-1

March 25, 2024

1 Ad Ease website Analytics

Dataset: Web Traffic Time Series Forecasting

Forecasting the future values of multiple time series. More specifically the problem of forecasting future web traffic for approximately 145,000 articles on britannica.

The training dataset consists of approximately 145k time series. Each of these time series represent a number of daily views of a different article, starting from July, 1st, 2015 up until December 31st, 2016. For each time series, you are provided the name of the article as well as the type of traffic that this time series represent (all, mobile, desktop, spider). You may use this metadata and any other publicly available data to make predictions. Unfortunately, the data source for this dataset does not distinguish between traffic values of zero and missing values. A missing value may mean the traffic was zero or that the data is not available for that day.

Data Dictionary:

there are two csv files given

train_1.csv :

In the csv file, each row corresponds to a particular article and each column correspond to a particular date. The values are the number of visits in that date.

Exog_Campaign_eng :

this file contains data for the dates which had a campaign or significant event that could affect the views for that day. the data is just for pages in english.

there is a 1 for dates with campaign and 0 for remaining dates. It is to be treated as an exogenous variable for models when training and forecasting data for pages in english

Importing the libraries

```
[ ]: !pip install --upgrade --no-cache-dir gdown
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
```

```
Requirement already satisfied: gdown in /usr/local/lib/python3.8/dist-packages (4.6.0)
```

```
Requirement already satisfied: six in /usr/local/lib/python3.8/dist-packages (from gdown) (1.15.0)
```

```
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.8/dist-
```

```

packages (from gdown) (2.23.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages
(from gdown) (4.64.1)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.8/dist-
packages (from gdown) (4.6.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.8/dist-
packages (from gdown) (3.8.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (2022.9.24)
Requirement already satisfied: urllib3!=1.25.0,!<1.25.1,<1.26,>=1.21.1 in
/usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (1.24.3)
Requirement already satisfied: chardet<4,>=3.0.2 in
/usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-
packages (from requests[socks]->gdown) (2.10)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in
/usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (1.7.1)

```

```

[ ]: import pandas as pd
import numpy as np
import pylab as p
import matplotlib.pyplot as plot
from collections import Counter
import re
import os
import seaborn as sns

```

```

[ ]: import warnings
warnings.filterwarnings("ignore")
warnings.simplefilter("ignore")

```

```

[ ]: sns.set(rc={'figure.figsize':(11.7,8.27)})

```

```

[ ]: import gdown
#url='https://drive.google.com/file/d/1gHYLqLt6rMyeAyvHf1wvLQ4BLKwjv9W/view?
↳usp=sharing'
#url='https://drive.google.com/file/d/1SL_7DoE16m71QpjJXoQUC3cI5aHCIZLv/view?
↳usp=share_link'
#url='https://drive.google.com/file/d/11GQSe2Xm4vFD4Xfw3JhOoPlXnBE_LiMe/view?
↳usp=sharing'
url='https://drive.google.com/file/d/1CJOMYyg64x3gN52p60qypN6UUgDnUhkm/view?
↳usp=sharing'

ider=url.split('/')[-2]
!gdown --id $ider

```

```

/usr/local/lib/python3.8/dist-packages/gdown/cli.py:121: FutureWarning: Option

```

`--id` was deprecated in version 4.3.1 and will be removed in 5.0. You don't need to pass it anymore to use a file ID.

```
warnings.warn(
```

```
Downloading...
```

```
From: https://drive.google.com/uc?id=1CJOMYyg64x3gN52p60qypN6UUgDnUhkm
```

```
To: /content/new_train.csv
```

```
100% 425M/425M [00:03<00:00, 125MB/s]
```

```
[ ]: train = pd.read_csv('new_train.csv')
```

Reading the dataset and printing head and tail to get basic idea

```
[ ]: train.head()
```

```
[ ]:
```

		Page	2015-07-01	2015-07-02	\
0	2NE1_zh.britanica.org_all-access_spider		18.0	11.0	
1	2PM_zh.britanica.org_all-access_spider		11.0	14.0	
2	3C_zh.britanica.org_all-access_spider		1.0	0.0	
3	4minute_zh.britanica.org_all-access_spider		35.0	13.0	
4	52_Hz_I_Love_You_zh.britanica.org_all-access_s...		NaN	NaN	

	2015-07-03	2015-07-04	2015-07-05	2015-07-06	2015-07-07	2015-07-08	\
0	5.0	13.0	14.0	9.0	9.0	22.0	
1	15.0	18.0	11.0	13.0	22.0	11.0	
2	1.0	1.0	0.0	4.0	0.0	3.0	
3	10.0	94.0	4.0	26.0	14.0	9.0	
4	NaN	NaN	NaN	NaN	NaN	NaN	

	2015-07-09	...	2016-12-22	2016-12-23	2016-12-24	2016-12-25	\
0	26.0	...	32.0	63.0	15.0	26.0	
1	10.0	...	17.0	42.0	28.0	15.0	
2	4.0	...	3.0	1.0	1.0	7.0	
3	11.0	...	32.0	10.0	26.0	27.0	
4	NaN	...	48.0	9.0	25.0	13.0	

	2016-12-26	2016-12-27	2016-12-28	2016-12-29	2016-12-30	2016-12-31
0	14.0	20.0	22.0	19.0	18.0	20.0
1	9.0	30.0	52.0	45.0	26.0	20.0
2	4.0	4.0	6.0	3.0	4.0	17.0
3	16.0	11.0	17.0	19.0	10.0	11.0
4	3.0	11.0	27.0	13.0	36.0	10.0

[5 rows x 551 columns]

```
[ ]: print(train.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 145063 entries, 0 to 145062
```

```
Columns: 551 entries, Page to 2016-12-31
dtypes: float64(550), object(1)
memory usage: 609.8+ MB
None
```


```
[ ]: print(train.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145063 entries, 0 to 145062
Columns: 551 entries, Page to 2016-12-31
dtypes: float64(550), object(1)
memory usage: 609.8+ MB
None
```

We can see that there are some null values in the data, we will plot them to see how it looks

```
[ ]: days = [r for r in range(1, len(train.columns))]
plot.figure(figsize=(10,7))
plot.xlabel('Day')
plot.ylabel('Null values')
plot.plot(days, train.isnull().sum()[1:])
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f9af53dca90>]
```



We see that the number of nan values decrease with time.

Probable reason: Some website have all nan values in the begining, that can be due to the fact that those were created after that time so there is no traffic reading for that time

```
[ ]: print(train.shape)
train=train.dropna(how='all')
# 'all' : If all values are NA, drop that row or column.
print(train.shape)

train=train.dropna(thresh=300)
print(train.shape)
```

(145063, 551)

(145063, 551)

(133617, 551)

1. We try dropping the rows that have all values as nan, none in our case.
2. We then also drop rows that have nan more than 300 days, because the time series for that would not make much sense
3. We fill all the remaining values with zero assuming there was no traffic on the date that the values are nan for.

```
[ ]: train=train.fillna(0)
train.tail()
```

```
[ ]:
                                     Page  2015-07-01  \
145012  Legión_(Marvel_Comics)_es.britanica.org_all-ac...      0.0
145013  Referéndum_sobre_la_permanencia_del_Reino_Unid...      0.0
145014  Salida_del_Reino_Unido_de_la_Unión_Europea_es...      0.0
145015  Amar,_después_de_amar_es.britanica.org_all-acc...      0.0
145016  Anexo:89.º_Premios_Óscar_es.britanica.org_all-...      0.0

      2015-07-02  2015-07-03  2015-07-04  2015-07-05  2015-07-06  \
145012          0.0          0.0          0.0          0.0          0.0
145013          0.0          0.0          0.0          0.0          0.0
145014          0.0          0.0          0.0          0.0          0.0
145015          0.0          0.0          0.0          0.0          0.0
145016          0.0          0.0          0.0          0.0          0.0

      2015-07-07  2015-07-08  2015-07-09  ...  2016-12-22  2016-12-23  \
145012          0.0          0.0          0.0  ...          7.0          3.0
145013          0.0          0.0          0.0  ...          9.0         16.0
145014          0.0          0.0          0.0  ...         29.0         36.0
145015          0.0          0.0          0.0  ...          7.0         30.0
145016          0.0          0.0          0.0  ...          0.0          0.0

      2016-12-24  2016-12-25  2016-12-26  2016-12-27  2016-12-28  \
```

145012	2.0	4.0	2.0	4.0	4.0
145013	8.0	3.0	6.0	3.0	3.0
145014	23.0	182.0	43.0	8.0	22.0
145015	27.0	14.0	8.0	7.0	5.0
145016	1.0	0.0	1.0	0.0	0.0

	2016-12-29	2016-12-30	2016-12-31
145012	1.0	2.0	2.0
145013	10.0	11.0	3.0
145014	13.0	18.0	14.0
145015	43.0	12.0	25.0
145016	1.0	1.0	0.0

[5 rows x 551 columns]

2 EDA

The page values are in this format

SPECIFIC NAME __ LANGUAGE.britanica.org __ ACCESS TYPE __ ACCESS ORIGIN

having information about page name, the main domain, device type used to access the page, and also the request origin(spider or browser agent)

```
[ ]: #Usage of Regex
def split_page(page):
    w = re.split('_|\.', page)
    print(w)
    return ' '.join(w[:-5]), w[-5], w[-2], w[-1]

split_page('2NE1_zh.britanica.org_all-access_spider')
```

```
['2NE1', 'zh', 'britanica', 'org', 'all-access', 'spider']
```

```
[ ]: ('2NE1', 'zh', 'all-access', 'spider')
```

```
[ ]: def split_page(page):
    w = re.split('_|\.', page)
    return ' '.join(w[:-5]), w[-5], w[-2], w[-1]

li = list(train.Page.apply(lambda x: split_page(str(x))))
df = pd.DataFrame(li)
df.columns = ['Title', 'Language', 'Access_type', 'Access_origin']
df = pd.concat([train, df], axis = 1)
```

We split the page name and get that information joining it with a temporary database. below we get some rows to see the structure of the data

```
[ ]: df.head()
```

```
[ ]:
      Page 2015-07-01 2015-07-02 \
0  2NE1_zh.britanica.org_all-access_spider 18.0 11.0
1  2PM_zh.britanica.org_all-access_spider 11.0 14.0
2  3C_zh.britanica.org_all-access_spider 1.0 0.0
3  4minute_zh.britanica.org_all-access_spider 35.0 13.0
4  NaN NaN NaN

      2015-07-03 2015-07-04 2015-07-05 2015-07-06 2015-07-07 2015-07-08 \
0  5.0 13.0 14.0 9.0 9.0 22.0
1  15.0 18.0 11.0 13.0 22.0 11.0
2  1.0 1.0 0.0 4.0 0.0 3.0
3  10.0 94.0 4.0 26.0 14.0 9.0
4  NaN NaN NaN NaN NaN NaN


      2015-07-09 ... 2016-12-26 2016-12-27 2016-12-28 2016-12-29 \
0  26.0 ... 14.0 20.0 22.0 19.0
1  10.0 ... 9.0 30.0 52.0 45.0
2  4.0 ... 4.0 4.0 6.0 3.0
3  11.0 ... 16.0 11.0 17.0 19.0
4  NaN ... NaN NaN NaN NaN

      2016-12-30 2016-12-31 Title Language Access_type Access_origin
0  18.0 20.0 2NE1 zh all-access spider
1  26.0 20.0 2PM zh all-access spider
2  4.0 17.0 3C zh all-access spider
3  10.0 11.0 4minute zh all-access spider
4  NaN NaN 5566 zh all-access spider
```

```
[5 rows x 555 columns]
```

```
[ ]: sns.countplot(df['Language'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9af603af10>
```




This above is the comparision number of articles in each language

{'ja':'Japanese', 'de':'German', 'en': 'English', 'no_lang':'Media_File', 'fr':'French', 'zh':'Chinese', 'ru':'Russian', 'es':'Spanish'}

```
[ ]: sns.countplot(df['Access_type'])
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9af249c220>
```

This comparison shows that usage from desktop and mobile is almost the same

```
[ ]: sns.countplot(df['Access_origin'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9af23f3880>
```



This shows that organic view is far more than that of spiders or bots

Now we want to compare the views for different languages

```
[ ]: #here we see that the languages are not treated properly as there are commons
      ↪and www
df.groupby('Language').count()
```

```
[ ]:
      Page  2015-07-01  2015-07-02  2015-07-03  2015-07-04  2015-07-05  \
Language
commons    7672         7672         7672         7672         7672
de         15946        15946        15946        15946        15946
en         20758        20758        20758        20758        20758
es         12268        12268        12268        12268        12268
fr         15418        15418        15418        15418        15418
ja         17132        17132        17132        17132        17132
ru         12955        12955        12955        12955        12955
www         5743         5743         5743         5743         5743
zh         14845        14845        14845        14845        14845

      2015-07-06  2015-07-07  2015-07-08  2015-07-09  ...  2016-12-25  \
Language
commons         7672         7672         7672         7672  ...         7672
de         15946        15946        15946        15946  ...        15946
```

en	20758	20758	20758	20758	...	20758
es	12268	12268	12268	12268	...	12268
fr	15418	15418	15418	15418	...	15418
ja	17132	17132	17132	17132	...	17132
ru	12955	12955	12955	12955	...	12955
www	5743	5743	5743	5743	...	5743
zh	14845	14845	14845	14845	...	14845

	2016-12-26	2016-12-27	2016-12-28	2016-12-29	2016-12-30	\
Language						
commons	7672	7672	7672	7672	7672	
de	15946	15946	15946	15946	15946	
en	20758	20758	20758	20758	20758	
es	12268	12268	12268	12268	12268	
fr	15418	15418	15418	15418	15418	
ja	17132	17132	17132	17132	17132	
ru	12955	12955	12955	12955	12955	
www	5743	5743	5743	5743	5743	
zh	14845	14845	14845	14845	14845	

	2016-12-31	Title	Access_type	Access_origin
Language				
commons	7672	8266	8266	8266
de	15946	17362	17362	17362
en	20758	22486	22486	22486
es	12268	13551	13551	13551
fr	15418	16948	16948	16948
ja	17132	19295	19295	19295
ru	12955	14270	14270	14270
www	5743	6228	6228	6228
zh	14845	15211	15211	15211

[9 rows x 554 columns]

```
[ ]: df[df['Language']=='commons']
```

```
[ ]:
                                     Page 2015-07-01 \
12271  Burning_Man_en.britanica.org_desktop_all-agents      1693.0
12272  Cali_Cartel_en.britanica.org_desktop_all-agents       348.0
12273  Call_of_Duty:_Modern_Warfare_2_en.britanica.or...     806.0
12274  Calvin_Harris_en.britanica.org_desktop_all-agents     7114.0
12275  Carl_Sagan_en.britanica.org_desktop_all-agents      1808.0
...
75274  Ash_Wednesday_en.britanica.org_mobile-web_all-...    170.0
75275  Ashley_Williams_(footballer)_en.britanica.org_...    112.0
75276  Assassin's_Creed_(film)_en.britanica.org_mobil...     28.0
75277  Aubrey_Plaza_en.britanica.org_mobile-web_all-a...   3067.0
```

75278 Australia_Plus_en.britanica.org_mobile-web_all... 17.0

	2015-07-02	2015-07-03	2015-07-04	2015-07-05	2015-07-06	2015-07-07 \
12271	1490.0	1186.0	1099.0	1051.0	1968.0	1874.0
12272	363.0	214.0	252.0	257.0	320.0	340.0
12273	768.0	700.0	725.0	723.0	823.0	849.0
12274	5599.0	7685.0	15844.0	9390.0	7173.0	5499.0
12275	1759.0	1838.0	1631.0	1701.0	2230.0	2052.0
...
75274	169.0	165.0	166.0	186.0	154.0	154.0
75275	102.0	135.0	147.0	120.0	89.0	90.0
75276	15.0	24.0	24.0	27.0	34.0	36.0
75277	2952.0	3459.0	3310.0	3294.0	3885.0	3830.0
75278	11.0	14.0	6.0	10.0	8.0	18.0

	2015-07-08	2015-07-09	...	2016-12-26	2016-12-27	2016-12-28 \
12271	1637.0	1842.0	...	940.0	2976.0	2447.0
12272	345.0	432.0	...	1197.0	1449.0	1376.0
12273	731.0	833.0	...	818.0	929.0	892.0
12274	6367.0	5630.0	...	1242.0	1409.0	1714.0
12275	2065.0	2002.0	...	1425.0	1730.0	1858.0
...
75274	184.0	173.0	...	475.0	444.0	389.0
75275	109.0	215.0	...	348.0	439.0	357.0
75276	35.0	22.0	...	43471.0	41147.0	38978.0
75277	3325.0	3150.0	...	4004.0	4848.0	5125.0
75278	10.0	6.0	...	12.0	17.0	18.0

	2016-12-29	2016-12-30	2016-12-31 \
12271	1648.0	1282.0	956.0
12272	1350.0	1287.0	992.0
12273	878.0	893.0	712.0
12274	1604.0	1602.0	1300.0
12275	1781.0	1718.0	1345.0
...
75274	382.0	362.0	393.0
75275	267.0	639.0	429.0
75276	36517.0	41760.0	38116.0
75277	4907.0	4617.0	4415.0
75278	13.0	6.0	19.0

	Title	Language \
12271	Accueil	commons
12272	Atlas of Asia	commons
12273	Atlas of Europe	commons
12274	Atlas of World War II	commons
12275	Atlas of colonialism	commons

```

...
75274 File:El jardín de las Delicias, de El Bosco jpg commons
75275 File:Jää on kulmunud pallideks (Looduse veidru... commons
75276 File:Reichstagsgebäude mit Weihnachtsbaum bei ... commons
75277 Category:Images by Eugène Cattin commons
75278 File:Nikolaos Gyzis - Historia jpg commons

```

```

Access_type Access_origin
12271 all-access spider
12272 all-access spider
12273 all-access spider
12274 all-access spider
12275 all-access spider

```

```

...
75274 desktop all-agents
75275 desktop all-agents
75276 desktop all-agents
75277 desktop all-agents
75278 desktop all-agents

```

[8266 rows x 555 columns]

```
[ ]: # Checking another way of fetching the language out of the string
```

```

def lang(Page):
    val = re.search('[a-z][a-z].britannica.org',Page)
    if val:
        #print(val)
        #print(val[0][0:2] )

        return val[0][0:2]

    return 'no_lang'

df['Language']=df['Page'].apply(lambda x: lang(str(x)))

```

```
[ ]: df.groupby('Language').count() #now the count has increased. You can go back
    ↪and get it sorted
```

```
[ ]:
      Page  2015-07-01  2015-07-02  2015-07-03  2015-07-04  2015-07-05  \
Language
de      17362      17362      17362      17362      17362      17362
en      22486      22486      22486      22486      22486      22486
es      13551      13551      13551      13551      13551      13551
fr      16948      16948      16948      16948      16948      16948
ja      19295      19295      19295      19295      19295      19295
no_lang 14494      14494      14494      14494      14494      14494
ru      14270      14270      14270      14270      14270      14270

```

zh	15211	15211	15211	15211	15211	15211
----	-------	-------	-------	-------	-------	-------

	2015-07-06	2015-07-07	2015-07-08	2015-07-09	...	2016-12-25 \
Language					...	
de	17362	17362	17362	17362	...	17362
en	22486	22486	22486	22486	...	22486
es	13551	13551	13551	13551	...	13551
fr	16948	16948	16948	16948	...	16948
ja	19295	19295	19295	19295	...	19295
no_lang	14494	14494	14494	14494	...	14494
ru	14270	14270	14270	14270	...	14270
zh	15211	15211	15211	15211	...	15211

	2016-12-26	2016-12-27	2016-12-28	2016-12-29	2016-12-30 \
Language					
de	17362	17362	17362	17362	17362
en	22486	22486	22486	22486	22486
es	13551	13551	13551	13551	13551
fr	16948	16948	16948	16948	16948
ja	19295	19295	19295	19295	19295
no_lang	14494	14494	14494	14494	14494
ru	14270	14270	14270	14270	14270
zh	15211	15211	15211	15211	15211

	2016-12-31	Title	Access_type	Access_origin
Language				
de	17362	13046	13046	13046
en	22486	22486	22486	22486
es	13551	10142	10142	10142
fr	16948	16948	16948	16948
ja	19295	16140	16140	16140
no_lang	14494	25374	25374	25374
ru	14270	14270	14270	14270
zh	15211	15211	15211	15211

[8 rows x 554 columns]

```
[ ]: df_language=df.groupby('Language').mean().transpose()
df_language
```

```
[ ]: Language      de      en      es      fr      ja \
2015-07-01  763.765926  3767.328604  1127.485204  499.092872  614.637160
2015-07-02  753.362861  3755.158765  1077.485425  502.297852  705.813216
2015-07-03  723.074415  3565.225696   990.895949  483.007553  637.451671
2015-07-04  663.537323  3711.782932   930.303151  516.275785  800.897435
2015-07-05  771.358657  3833.433025  1011.759575  506.871666  768.352319
...          ...          ...          ...          ...          ...
```

2016-12-27	1119.596936	6314.335275	1070.923400	840.590217	808.541436
2016-12-28	1062.284069	6108.874144	1108.996753	783.585379	807.430163
2016-12-29	1033.939062	6518.058525	1058.660320	763.209169	883.752786
2016-12-30	981.786430	5401.792360	807.551177	710.502773	979.278777
2016-12-31	937.842875	5280.643467	776.934322	654.060656	1228.720808


Language	no_lang	ru	zh
2015-07-01	102.733545	663.199229	272.498521
2015-07-02	107.663447	674.677015	272.906778
2015-07-03	101.769629	625.329783	271.097167
2015-07-04	86.853871	588.171829	273.712379
2015-07-05	96.254105	626.385354	291.977713
...
2016-12-27	155.270181	998.374071	363.066991
2016-12-28	178.561267	945.054730	369.049701
2016-12-29	150.873534	909.352207	340.526330
2016-12-30	156.049193	815.475123	342.745316
2016-12-31	135.792052	902.600210	352.184275

[550 rows x 8 columns]

```
[ ]: df_language.reset_index(inplace=True)
df_language.set_index('index', inplace=True)
```

```
[ ]: df_language.plot(figsize=(12,7))
plot.ylabel('Views per Page')
```

```
[ ]: Text(0, 0.5, 'Views per Page')
```




Plotting the data shows that articles in english get the most number of views as compared to different languages, there are some spikes at different times in different languages

Plotting just for english because we are going to use this for our further investigation and predictions

```
[ ]: df_language['en'].plot(figsize=(12,7))
      plot.ylabel('Views per Page')
```

```
[ ]: Text(0, 0.5, 'Views per Page')
```



```
[ ]: total_view=df_language.copy()
```

```
[ ]: #####
```

3 Checking the stationarity

Dickey-Fuller test

Here the null hypothesis is that the TS is non-stationary: The test results comprise of a Test Statistic and some Critical Values for difference confidence levels.

```
[ ]: from statsmodels.tsa.stattools import adfuller
      def df_test(x):
          result=adfuller(x)
```



```

print('ADF Stastistic: %f'%result[0])
print('p-value: %f'%result[1])

df_test(total_view['en'])

```

ADF Stastistic: -2.373563
p-value: 0.149337

We see that the p value is not low enough (<0.05). Therefore, we can say our series is not stationary as we fail to reject the null hypothesis

4 Making the time series stationary

```
[ ]: ts=total_view['en']
```

4.1 1. Remove trend and seasonality with decomposition

```
[ ]: # Naive decomposition of our Time Series as explained above
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(ts.values, model='multiplicative', freq = 7)
```

*""" Additive or multiplicative?
It's important to understand what the difference between a multiplicative
↪time series and an additive one before we go any further.*

There are three components to a time series:

- trend how things are overall changing
- seasonality how things change within a given period e.g. a year, month,
↪week, day
- error/residual/irregular activity not explained by the trend or the
↪seasonal value

*How these three components interact determines the difference between a
↪multiplicative and an additive time series.*

*In a multiplicative time series, the components multiply together to make the
↪time series. If you have an increasing trend, the amplitude of seasonal
↪activity increases. Everything becomes more exaggerated. This is common when
↪you're looking at web traffic.*

*In an additive time series, the components add together to make the time
↪series. If you have an increasing trend, you still see roughly the same size
↪peaks and troughs throughout the time series. This is often seen in indexed
↪time series where the absolute value is growing but changes stay relative.*


```

"""

trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plot.figure(figsize=(10,7))
plot.subplot(411)
plot.title('Observed = Trend + Seasonality + Residuals')
plot.plot(ts.values,label='Observed')
plot.legend(loc='best')
plot.subplot(412)
plot.plot(trend, label='Trend')
plot.legend(loc='best')
plot.subplot(413)
plot.plot(seasonal,label='Seasonality')
plot.legend(loc='best')
plot.subplot(414)
plot.plot(residual, label='Residuals')
plot.legend(loc='best')
plot.tight_layout()
plot.show()

```



```
[ ]: ts_decompose=pd.DataFrame(residual).fillna(0)[0].values
df_test(ts_decompose)
```


ADF Stastistic: -3.796320

p-value: 0.002945

We can see that our series is now stationary, we can also try differencing to see what results we can get.

5 2. Remove trend and seasonality with differencing

```
[ ]: ts_diff = ts - ts.shift(1)
plot.plot(ts_diff.values)
plot.show()
```



```
[ ]: ts_diff.dropna(inplace=True)
df_test(ts_diff)
```

ADF Stastistic: -8.273590

p-value: 0.000000

Also the p value is 0. So we can say that our graph is now stationery. Now we can apply the ARIMA model

How do we choose p,d,q


a thumb rule that for choosing the p,q values are when the lag goes below the significant level - we use PACF for p, here we see that till lag 5 there are significant lines, if we want our model to be simpler we can start with a smaller number like 3/4 - we use ACF for q. here we can see that lag 4 is below significant level so we will use till lag 3


as for d we can see that at 1 differencing the series becomes stationary so we choose d as 1

6 Plot the autocorreltaion and partial auto correlation functions

Plotting the graphs and getting the p,q,d values for arima

```
[ ]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
acf=plot_acf(ts_diff,lags=20)
pacf=plot_pacf(ts_diff,lags=20)
```





<https://people.duke.edu/~rnau/411arim3.htm>

[]:

7 ARIMA MODEL

```
[ ]: from statsmodels.tsa.arima_model import ARIMA
     from pandas import DataFrame
```

```
[ ]: model = ARIMA(ts, order=(4,1,3))
     model_fit = model.fit(dis=0)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.
```

```
warnings.warn('No frequency information was'
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.
```

```
warnings.warn('No frequency information was'
/usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:547:
HessianInversionWarning: Inverting hessian failed, no bse or cov_params
```

available


```
warnings.warn('Inverting hessian failed, no bse or cov_params '
/usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:566:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
warnings.warn("Maximum Likelihood optimization failed to "
```

```
[ ]: model_fit.plot_predict(dynamic=False)
"""When you set dynamic=True, the model continuously predicts one-step ahead,
    ↳ (t+1) and then for the 2nd step ahead (t+2) prediction, it appends predicted
    ↳ value (t+1) to data, re-fits model on new expanded data then makes 2nd step
    ↳ ahead forecast. This is called out-of-sample prediction.

    When you set dynamic=False, the model sequentially predicts one-step-ahead,
    ↳ using the true value from previous time step instead of using predicted
    ↳ value. This is called in-sample prediction.

    On your first comparison of plots as you predict from 509 to 533, the reason
    ↳ you get same plots is you are extrapolating, you do not have true values of
    ↳ next 24 steps that you predicted therefor regardless of setting dynamic
    ↳ either True or False model uses out-of-sample approach.

    Since out-of-sample approach uses the last predicted value from the previous
    ↳ time step to predict the next value in time, as number of steps get farther,
    ↳ it is expected to deviate from actual values because on each step's
    ↳ prediction fitted model learns previous predicted step's errors as well.
    """
plot.show()
```



[]:

Multistep forecasting

```
[ ]: train = ts[:-20]
     test = ts[-20:]
```

```
[ ]: model = ARIMA(train, order=(4, 1, 3))
     fitted = model.fit(dispatch=-1)

     # Forecast
     fc, se, conf = fitted.forecast(20, alpha=0.02)

     # Make as pandas series
     fc_series = pd.Series(fc, index=test.index)
     # Plot
     plot.figure(figsize=(12,5), dpi=100)
     plot.plot(train, label='training')
     plot.plot(test, label='actual')
     plot.plot(fc_series, label='forecast')

     plot.title('Forecast vs Actuals')
```

```
plot.legend(loc='upper left', fontsize=8)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.
```

```
warnings.warn('No frequency information was'
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.
```

```
warnings.warn('No frequency information was'
```


```
/usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:547:
HessianInversionWarning: Inverting hessian failed, no bse or cov_params
available
```

```
warnings.warn('Inverting hessian failed, no bse or cov_params '
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:566:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
[ ]: <matplotlib.legend.Legend at 0x7f9aed53c340>
```



```
[ ]: mape = np.mean(np.abs(fc - test.values)/np.abs(test.values))
rmse = np.mean((fc - test.values)**2)**.5
print("mape:", mape)
print("rsme:", rmse)
```

```
mape: 0.0935492560101219
```

```
rsme: 706.602453194038
```

we can see that the model does not perform very well for multistep out sample data

from the decomposition we can see that there is a weekly seasonality and still some spikes in the

residual, that may be because of some external factors, which we can take into account by using them as our exogenous variable

```
[ ]: !gdown 1H9054-eVP9IdANPOblXwX7Nd2r_Sjf1u
```

Downloading...

From: https://drive.google.com/uc?id=1H9054-eVP9IdANPOblXwX7Nd2r_Sjf1u

To: /content/Exog_Campaign_eng

100% 1.10k/1.10k [00:00<00:00, 1.79MB/s]

```
[ ]: ex_df = pd.read_csv('Exog_Campaign_eng')
     ex_df.head()
```

```
[ ]:      Exog
     0      0
     1      0
     2      0
     3      0
     4      0
```

We get the exogenous data from this csv file for english pages

```
[ ]: exog=ex_df['Exog'].to_numpy()
```

we will train a sarimax model for that and see if we get any improvements from using the two information.

the seasonal order and the values of PDQ are based upon various trials and comparison of the models - we see a seasonality of 7 from the plots ie: weekly seasonality (from the plots we can see that after some insignificant plots we have some significant values repeating at intervals of 7 ie: 7,14 ...) - the non seasonal order we can keep the same

```
[ ]: import statsmodels.api as sm
     train=ts[:520]
     test=ts[520:]
     model=sm.tsa.statespace.SARIMAX(train,order=(4, 1, 1,
     ↪3),seasonal_order=(1,1,1,7),exog=exog[:520])
     results=model.fit()

     fc=results.forecast(30,dynamic=True,exog=pd.DataFrame(exog[520:]))

     # Make as pandas series
     fc_series = pd.Series(fc)
     # Plot
     train.index=train.index.astype('datetime64[ns]')
     test.index=test.index.astype('datetime64[ns]')
     plot.figure(figsize=(12,5), dpi=100)
     plot.plot(train, label='training')
     plot.plot(test, label='actual')
```

```

plot.plot(fc_series, label='forecast')

plot.title('Forecast vs Actuals')
plot.legend(loc='upper left', fontsize=8)

```

```

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.

```

```

warnings.warn('No frequency information was')
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524:
ValueWarning: No frequency information was provided, so inferred frequency D
will be used.

```

```

warnings.warn('No frequency information was')
/usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:566:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals

```

```


warnings.warn("Maximum Likelihood optimization failed to ")

```

```

[ ]: <matplotlib.legend.Legend at 0x7f9ae2bd5dc0>

```



```

[ ]:

```

```

[ ]: mape = np.mean(np.abs(fc - test.values)/np.abs(test.values))
rmse = np.mean((fc - test.values)**2)**.5
print("mape:", mape)
print("rsme:", rmse)

```

```

mape: 0.0476009066291969
rsme: 299.17343793278815

```

The mean absolute percentage error and the root mean squared error is low

8 regression for a time series

```
[ ]: ts_df=ts.to_frame()  
ts_df.head()
```

```
[ ]:          en  
index  
2015-07-01  3767.328604  
2015-07-02  3755.158765  
2015-07-03  3565.225696  
2015-07-04  3711.782932  
2015-07-05  3833.433025
```

```
[ ]: ts_df.reset_index(level=0, inplace=True)  
ts_df['date']=pd.to_datetime(ts_df['index'])  
ts_df.drop(['index'],axis=1,inplace=True)  
ts_df.head()
```

```
[ ]:          en      date  
0  3767.328604 2015-07-01  
1  3755.158765 2015-07-02  
2  3565.225696 2015-07-03  
3  3711.782932 2015-07-04  
4  3833.433025 2015-07-05
```

```
[ ]: ts_df['day_of_week']=ts_df['date'].dt.day_name()  
ts_df.head()
```

```
[ ]:          en      date day_of_week  
0  3767.328604 2015-07-01   Wednesday  
1  3755.158765 2015-07-02    Thursday  
2  3565.225696 2015-07-03     Friday  
3  3711.782932 2015-07-04    Saturday  
4  3833.433025 2015-07-05     Sunday
```

```
[ ]: ts_df=pd.get_dummies(ts_df, columns = ['day_of_week'])
```

```
[ ]: ts_df.head()
```

```
[ ]:          en      date  day_of_week_Friday  day_of_week_Monday  \  
0  3767.328604 2015-07-01                0                0  
1  3755.158765 2015-07-02                0                0  
2  3565.225696 2015-07-03                1                0  
3  3711.782932 2015-07-04                0                0  
4  3833.433025 2015-07-05                0                0  
  
    day_of_week_Saturday  day_of_week_Sunday  day_of_week_Thursday  \  

```

0	0	0	0
1	0	0	1
2	0	0	0
3	1	0	0
4	0	1	0

	day_of_week_Tuesday	day_of_week_Wednesday
0	0	1
1	0	0
2	0	0
3	0	0
4	0	0

```
[ ]: ts_df['exog']=ex_df['Exog']
ts_df['rolling_mean']=ts_df['en'].rolling(7).mean()
```

```
[ ]:
```

```
[ ]: ts_df=ts_df.dropna()
ts_df.head()
```

```
[ ]:
      en      date  day_of_week_Friday  day_of_week_Monday  \
6   3906.341724  2015-07-07              0              0
7   3685.854621  2015-07-08              0              0
8   3771.183714  2015-07-09              0              0
9   3749.860313  2015-07-10              1              0
10  3770.749355  2015-07-11              0              0
```

	day_of_week_Saturday	day_of_week_Sunday	day_of_week_Thursday	\
6	0	0	0	
7	0	0	0	
8	0	0	1	
9	0	0	0	
10	1	0	0	

	day_of_week_Tuesday	day_of_week_Wednesday	exog	rolling_mean
6	1	0	0	3809.528545
7	0	1	0	3797.889404
8	0	0	0	3800.178683
9	0	0	0	3826.555056
10	0	0	0	3834.978831

```
[ ]: X=ts_df[['day_of_week_Friday',      'day_of_week_Monday',      'day_of_week_Saturday',
      ↪copy()
y=ts_df[['en']]

train_x = X[:-20]
```

```
test_x = X[-20:]

train_y = y[:-20]
test_y = y[-20:]
```

```
[ ]: from sklearn.linear_model import LinearRegression

# Train and pred
model = LinearRegression()
model.fit(train_x, train_y)
y_pred = (model.predict(test_x))

mape = np.mean(np.abs(y_pred - test_y.values)/np.abs(test_y.values))
print("mape:", mape)
```

```
mape: 0.04523968736329716
```

We can see here that our mape is better than our arima model but worse than our sarimax model

- Linear Regression Is Limited to Linear Relationships and in our case there is not a lot of linear relationship.
- it would have been better to use a regression based model for forecasting if we can build some better features.
- we have our series data and the exogenous variables, we add the day of week feature, other than that there are not a lot of features that we can build

9 using Facebook Prophet

```
[ ]: !pip install pystan~=2.14
!pip install fbprophet
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Requirement already satisfied: pystan~=2.14 in /usr/local/lib/python3.7/dist-
packages (2.19.1.1)
Requirement already satisfied: Cython!=0.25.1,>=0.22 in
/usr/local/lib/python3.7/dist-packages (from pystan~=2.14) (0.29.32)
Requirement already satisfied: numpy>=1.7 in /usr/local/lib/python3.7/dist-
packages (from pystan~=2.14) (1.21.6)
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Requirement already satisfied: fbprophet in /usr/local/lib/python3.7/dist-
packages (0.7.1)
Requirement already satisfied: holidays>=0.10.2 in
/usr/local/lib/python3.7/dist-packages (from fbprophet) (0.14.2)
Requirement already satisfied: cmdstanpy==0.9.5 in
/usr/local/lib/python3.7/dist-packages (from fbprophet) (0.9.5)
```

Requirement already satisfied: convertdate>=2.1.2 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (2.4.0)

Requirement already satisfied: matplotlib>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (3.2.2)

Requirement already satisfied: setuptools-git>=1.2 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (1.2)

Requirement already satisfied: numpy>=1.15.4 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (1.21.6)

Requirement already satisfied: tqdm>=4.36.1 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (4.64.0)

Requirement already satisfied: pystan>=2.14 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (2.19.1.1)

Requirement already satisfied: python-dateutil>=2.8.0 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (2.8.2)

Requirement already satisfied: Cython>=0.22 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (0.29.32)

Requirement already satisfied: LunarCalendar>=0.0.9 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (0.0.9)

Requirement already satisfied: pandas>=1.0.4 in /usr/local/lib/python3.7/dist-packages (from fbprophet) (1.3.5)

Requirement already satisfied: pymeeus<=1,>=0.3.13 in /usr/local/lib/python3.7/dist-packages (from convertdate>=2.1.2->fbprophet) (0.5.11)

Requirement already satisfied: hijri-converter in /usr/local/lib/python3.7/dist-packages (from holidays>=0.10.2->fbprophet) (2.2.4)

Requirement already satisfied: korean-lunar-calendar in /usr/local/lib/python3.7/dist-packages (from holidays>=0.10.2->fbprophet) (0.2.1)

Requirement already satisfied: pytz in /usr/local/lib/python3.7/dist-packages (from LunarCalendar>=0.0.9->fbprophet) (2022.2.1)

Requirement already satisfied: ephemeris>=3.7.5.3 in /usr/local/lib/python3.7/dist-packages (from LunarCalendar>=0.0.9->fbprophet) (4.1.3)

Requirement already satisfied: cyclical>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.0.0->fbprophet) (0.11.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.0.0->fbprophet) (1.4.4)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.0.0->fbprophet) (3.0.9)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1->matplotlib>=2.0.0->fbprophet) (4.1.1)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.8.0->fbprophet) (1.15.0)

```
[ ]: ts_df['ds']=ts_df['date']
      ts_df['y']=ts_df['en']
```

```
[ ]: df2=ts_df[['date','en','exog']].copy()
      df2.columns = ['ds', 'y', 'exog']
      df2.head()
```

```
[ ]:
      ds          y  exog
6  2015-07-07  3906.341724    0
7  2015-07-08  3685.854621    0
8  2015-07-09  3771.183714    0
9  2015-07-10  3749.860313    0
10 2015-07-11  3770.749355    0
```

```
[ ]: df2[:20].info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 524 entries, 6 to 529
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    ds      524 non-null    datetime64[ns]
1    y        524 non-null    float64
2    exog     524 non-null    int64
dtypes: datetime64[ns](1), float64(1), int64(1)
memory usage: 16.4 KB
```

prophet without exogenous

```
[ ]: from fbprophet import Prophet
      m = Prophet(weekly_seasonality=True)
      m.fit(df2[['ds', 'y']][:20])
      future = m.make_future_dataframe(periods=20,freq="D")
      forecast = m.predict(future)
      fig = m.plot(forecast)
```

INFO:fbprophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to override this.

INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.




prophet with exogenous

```
[ ]: model2=Prophet(interval_width=0.9, weekly_seasonality=True,
    ↪changeoint_prior_scale=1)
model2.add_regressor('exog')
model2.fit(df2[:20])
forecast2 = model2.predict(df2)
fig = model2.plot(forecast2)
```


INFO:fbprophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to override this.

INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.



```
[ ]: y_true = df2['y'].values
      y_pred = forecast2['yhat'].values

      plot.plot(y_true, label='Actual')
      plot.plot(y_pred, label='Predicted')
      plot.legend()
      plot.show()
```



```
[ ]: mape = np.mean(np.abs(forecast2['yhat'][-20:] - df2['y'][-20:].values)/np.
      ↪abs(df2['y'][-20:].values))
      print("mape:",mape)
```

mape: 0.06592815614410931

- Prophet does not perform well on non-stationary data because it is difficult to find the actual seasonality and trend of the data if the patterns are inconsistent.

10

11 Comparing the predicted views for different languages

For doing this we are going to automate the procedure from loading the separate data for each language to doing out of sample forecasting for the next month, and then comparing the results.

```
[ ]: def grid_search(ts):
      v=[0,1,2,3]
      mape=100
      val=[0,0,0]
      for p in v:
```

```

        for d in v:
            for q in v:
                try:
                    model = ARIMA(ts[:-20], order=(p,d,q))
                    model_fit = model.fit(dispatch=-1)
                    fc, se, conf = model_fit.forecast(20, alpha=0.02)
                    x = np.mean(np.abs(fc - ts[-20:].values)/np.abs(ts[-20:].
↪values))

                    if(x<mape):
                        mape=x
                        val=[p,d,q]

                except:
                    pass

    return(mape, val)

```

This functions works like a grid search for getting the best value of p,d,q by comparing the mape of all models that we create.

the values of p,d,q that give the least mape score are saved and returned

```

[ ]: def all_arima(train,test,val):
    model = ARIMA(train, order=(val[0], val[1], val[2]))
    fitted = model.fit(dispatch=-1)

    # Forecast
    fc, se, conf = fitted.forecast(30, alpha=0.02)

    fc_series = pd.Series(fc, index=test.index)

    # Plot
    plot.figure(figsize=(12,5), dpi=100)
    plot.plot(train, label='training')
    plot.plot(test, label='actual')
    plot.plot(fc_series, label='forecast')
    plot.title('Forecast vs Actuals')
    plot.legend(loc='upper left', fontsize=8)
    plot.show()
    mape = np.mean(np.abs(fc - test.values)/np.abs(test.values))
    rmse = np.mean((fc - test.values)**2)**.5
    print("mape:",mape)
    print("rsme:",rmse)
    return (fc)

```


This function takes the p,d,q values that we calculated earlier and then trains a model on it, does

forecast and plots them for visualization.


it also calculates the sum of forecasted views for the next 30 days and returns it back

```
[ ]: import warnings
warnings.filterwarnings("ignore")
views_prediction={}
for c in total_view:
    print("language: ",c)
    ts=(total_view[c])
    mape,val=grid_search(ts)
    print(mape,val)
    train = ts[:520]
    test = ts[520:]
    fc=all_arima(train,test,val)
    views_prediction[c]=fc
```


```
language:  de
0.09397758421308047 [3, 1, 3]
```




```
mape: 0.08451930259659801
rsme: 119.84846155914948
language:  en
0.07497050608151384 [3, 2, 0]
```




mape: 0.4781050962988165
rsme: 2972.249073736557
language: es
0.08748369803155673 [3, 1, 3]




mape: 0.3032163163374593
rsme: 308.3268548323757
language: fr
0.07922539990107266 [0, 0, 2]




mape: 0.07316960356482773
rsme: 82.05815499130937
language: ja
0.09306005747257282 [0, 0, 0]




mape: 0.08086798857229117
rsme: 111.15944936830843
language: no_lang
0.07911978906564705 [0, 0, 3]



mape: 0.08096654200730383
rsme: 21.960681343449338
language: ru
0.052023758161538015 [2, 2, 3]



mape: 0.1196084671042122
rsme: 122.51049610367977
language: zh
0.04467696433881181 [0, 0, 0]



mape: 0.04704826775436991

rsme: 20.17502184263151

##Difference between arima, sarima & sarimax.

ARIMA (AutoRegressive Integrated Moving Average), SARIMA (Seasonal AutoRegressive Integrated Moving Average), and SARIMAX (Seasonal AutoRegressive Integrated Moving Average with eXogenous factors) are all models used in time series analysis for forecasting. They are extensions or variations of the basic ARIMA model, incorporating different features to handle specific characteristics of time series data.

1. ARIMA (AutoRegressive Integrated Moving Average):

- ARIMA models are used for non-seasonal time series data.
- It combines three components: AutoRegressive (AR), Integrated (I), and Moving Average (MA).
- AR component captures the relationship between an observation and a lagged observation in the same time series.
- I component involves differencing the time series to make it stationary, i.e., to remove trends or seasonality.
- MA component represents the relationship between an observation and a residual error from a moving average model applied to lagged observations.

2. SARIMA (Seasonal AutoRegressive Integrated Moving Average):

- SARIMA models extend ARIMA to handle seasonal variations in time series data.
- It includes additional seasonal AR, seasonal differencing, and seasonal MA components.
- The seasonal AR component captures the relationship between an observation and a lagged observation from the same season.
- Seasonal differencing is performed to remove seasonal patterns.
- The seasonal MA component models the relationship between an observation and a residual error from a seasonal moving average model applied to lagged observations.

3. SARIMAX (Seasonal AutoRegressive Integrated Moving Average with eXogenous factors):

- SARIMAX extends SARIMA by allowing the inclusion of exogenous variables, i.e., ex-

ternal factors that can influence the time series.

- In addition to the seasonal ARIMA components, SARIMAX incorporates exogenous variables into the model.
- Exogenous variables are independent variables that are not influenced by the time series itself but may have an impact on its behavior.
- Including exogenous variables in the model can improve forecast accuracy by accounting for additional factors that may affect the time series.

In summary, ARIMA is used for non-seasonal time series, SARIMA extends ARIMA to handle seasonal patterns, and SARIMAX further extends SARIMA by allowing the inclusion of exogenous variables to better capture the complexities of real-world time series data.